



ACC-HPC June 2025

UNSUPERVISED LEARNING & CLUSTERING

Dr Kiran Waghmare
CDAC Mumbai

Machine Learning >




Agenda

- Clustering
- K-Means
- Hierarchical
- DB-SCAN

Machine learning:

- **Supervised vs Unsupervised.**
 - **Supervised learning** - the presence of the outcome variable is available to guide the learning process.
 - there must be a training data set in which the solution is already known.
 - **Unsupervised learning** - the outcomes are unknown.
 - cluster the data to reveal meaningful partitions and hierarchies

Unsupervised Learning



Clustering

Clustering:

- Unsupervised learning
- Requires data, but no labels
- Detect patterns e.g. in
 - Group emails or search results
 - Customer shopping patterns
 - Regions of images
- Useful when don't know what you're looking for
- But: can get gibberish



Machine Learning: Clustering



By color



By shape



By size



etc...

Cluster by Type

Clustering:

- Clustering is the task of gathering samples into groups of similar samples according to some predefined similarity or dissimilarity measure





sample




Cluster/group

- In this case clustering is carried out using the Euclidean distance as a measure.



Clustering

- Basic idea: group together similar instances
- Example: 2D point patterns



What is Similarity?

Slide based on one by Eamonn Keogh



Similarity is
hard to define,
but...
*“We know it
when we see it”*

Partitioning Algorithms: Basic Concept


- Partitioning method: Construct a partition of a database D of n objects into a set of k clusters, s.t., min sum of squared distance

$$\sum_{m=1}^k \sum_{t_{mi} \in K_m} (C_m - t_{mi})^2$$

- Given a k , find a partition of k clusters that optimizes the chosen partitioning criterion
 - Global optimal: exhaustively enumerate all partitions
 - Heuristic methods: *k-means* and *k-medoids* algorithms
 - *k-means* (MacQueen'67): Each cluster is represented by the center of the cluster
 - *k-medoids* or PAM (Partition around medoids) (Kaufman & Rousseeuw'87): Each cluster is represented by one of the objects in the cluster


The *K-Means* Clustering Method

- Given k , the *k-means* algorithm is implemented in four steps:
 - Partition objects into k nonempty subsets
 - Compute seed points as the centroids of the clusters of the current partition (the centroid is the center, i.e., *mean point*, of the cluster)
 - Assign each object to the cluster with the nearest seed point
 - Go back to Step 2, stop when no more new assignment





The *K*-Means Clustering Method

■ Example




K=2
Arbitrarily choose K object as initial cluster center


↑
Assign each objects to most similar center



→
Update the cluster means




←
Update the cluster means



↓
reassign


K-Means

- An iterative clustering algorithm
 - Initialize: Pick K random points as cluster centers
 - Alternate:
 1. Assign data points to closest cluster center
 2. Change the cluster center to the average of its assigned points
 - Stop when no points' assignments change




K-Means

- An iterative clustering algorithm
 - Initialize: Pick K random points as cluster centers
 - Alternate:
 1. Assign data points to closest cluster center
 2. Change the cluster center to the average of its assigned points
 - Stop when no points' assignments change




K-means clustering: Example



- Pick K random points as cluster centers (means)

Shown here for $K=2$


K-means clustering: Example



Iterative Step 1

- Assign data points to closest cluster center


K-means clustering: Example



Iterative Step 2


- Change the cluster center to the average of the assigned points

K-means clustering: Example




- Repeat until convergence


K-means clustering: Example



K-means clustering: Example



K-means clustering: Example




Comments on the *K-Means* Method

- Strength: Relatively efficient: $O(tkn)$, where n is # objects, k is # clusters, and t is # iterations. Normally, $k, t \ll n$.
 - Comparing: PAM: $O(k(n-k)^2)$, CLARA: $O(ks^2 + k(n-k))$
- Comment: Often terminates at a *local optimum*. The *global optimum* may be found using techniques such as: *deterministic annealing* and *genetic algorithms*
- Weakness
 - Applicable only when *mean* is defined, then what about categorical data?
 - Need to specify k , the *number* of clusters, in advance
 - Unable to handle noisy data and *outliers*
 - Not suitable to discover clusters with *non-convex shapes*


High-Dimensional Data poses Problems for Clustering

- **Difficult to find true clusters**
 - Irrelevant and redundant features
 - All points are equally close
- **Solutions: Dimension Reduction**
 - Feature subset selection
 - Cluster ensembles using random projection (in a later lecture....)

Redundant



Irrelevant



Curse of Dimensionality


**100 observations cover the 1-D unit interval
[0,1] well**

**Consider the 10-D unit hypersquare, 100
observations are now isolated points in a
vast empty space.**

Consequence of the Curse

- Suppose the number of samples given to us in the total sample space is fixed
- Let the dimension increase
- Then the distance of the k nearest neighbors of any point increases

Machine Learning Fundamentals...




...Bias and Variance!!!

- In previous lectures we showed how to build classifiers when the underlying densities are known
 - Bayesian Decision Theory introduced the general formulation
- In most situations, however, the true distributions are unknown and must be estimated from data.
 - **Parameter Estimation** (we saw the Maximum Likelihood Method)
 - Assume a particular form for the density (e.g. Gaussian), so only the parameters (e.g., mean and variance) need to be estimated
 - Maximum Likelihood
 - Bayesian Estimation
 - **Non-parametric Density Estimation** (not covered)
 - Assume NO knowledge about the density
 - Kernel Density Estimation
 - Nearest Neighbor Rule

Bias and variance (1)


- How good are these estimates? Two measures of “goodness” are used for statistical estimates


- **BIAS:** how close is the estimate to the true value?
- **VARIANCE:** how much does the estimate change for different runs (e.g. different datasets)?



- The bias-variance tradeoff


- In most cases, you can only decrease one of them at the expense of the other








Underfitting

Overfitting







Principal component analysis (PCA)

- Widely used method for unsupervised, linear dimensionality reduction
- GOAL: account for variance of data in as few dimensions as possible (using linear projection)


Geometric picture of principal components (PCs)

- First PC is the projection direction that maximizes the variance of the projected data
- Second PC is the projection direction that is orthogonal to the first PC and maximizes variance of the projected data




PCA: conceptual algorithm

- Find a line, such that when the data is projected onto that line, it has the maximum variance.




PCA: conceptual algorithm


- Find a second line, orthogonal to the first, that has maximum projected variance.



PCA: conceptual algorithm


- Repeat until have k orthogonal lines
- The projected position of a point on these lines gives the coordinates in the k -dimensional reduced space.






PCA in a nutshell

1. correlated hi-d data
("urefu" means "height" in Swahili)




2. center the points




want dimension of highest variance

7. uncorrelated low-d data



6. project data points to those eigenvectors



3. compute covariance matrix

$$\begin{matrix} h & u \\ h & 2.0 & 0.8 \\ u & 0.8 & 0.6 \end{matrix} \rightarrow \text{cov}(h, u) = \frac{1}{n} \sum_{i=1}^n h_i u_i$$


4. eigenvectors + eigenvalues

$$\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} e_h \\ e_u \end{pmatrix} = \Lambda_e \begin{pmatrix} e_h \\ e_u \end{pmatrix}$$

$$\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} f_h \\ f_u \end{pmatrix} = \Lambda_f \begin{pmatrix} f_h \\ f_u \end{pmatrix}$$

`eig(cov(data))`

5. pick m<d eigenvectors w. highest eigenvalues



Steps in principal component analysis

- Mean center the data
- Compute covariance matrix Σ
- Calculate eigenvalues and eigenvectors of Σ
 - Eigenvector with largest eigenvalue λ_1 is 1st principal component (PC)
 - Eigenvector with k^{th} largest eigenvalue λ_k is k^{th} PC
 - $\lambda_k / \sum_i \lambda_i$ = proportion of variance captured by k^{th} PC

Applying a principal component analysis

- Full set of PCs comprise a new orthogonal basis for feature space, whose axes are aligned with the maximum variances of original data.
- Projection of original data onto first k PCs gives a reduced dimensionality representation of the data.
- Transforming reduced dimensionality projection back into original space gives a reduced dimensionality *reconstruction* of the original data.
- Reconstruction will have some error, but it can be small and often is acceptable given the other benefits of dimensionality reduction.